

COMS 4721: Machine Learning for Data Science

Columbia University, Spring 2015

Homework 2: Due March 3, 2015

Submit the written portion of your homework as a single PDF file through Courseworks (less than 5MB). In addition to your PDF write-up, submit all code written by you in their original extensions through Courseworks. Do not submit in .rar, .zip, .tar, .doc, or other file types. Your grade will be based on the contents of one PDF file and original source code. Everything resulting from the problems on this homework other than the raw code should be put in the PDF file.

Show all work for full credit. Late homeworks will not be accepted – i.e., homework submitted to Courseworks after midnight on the due date.

Problem 1 (multiclass logistic regression) – 15 points

Logistic regression with more than two classes can be done using the softmax function. For data $x \in \mathbb{R}^d$ and k classes (where class i has regression vector w_i) the class of x , denoted by y , follows the probability distribution

$$P(y|x, w_1, \dots, w_k) = \prod_{i=1}^k \left(\frac{e^{x^T w_i}}{\sum_{j=1}^k e^{x^T w_j}} \right)^{\mathbb{1}(y=i)}.$$

1. Write out the log likelihood \mathcal{L} of data $(x_1, y_1), \dots, (x_n, y_n)$ using an i.i.d. assumption.
2. Calculate $\nabla_{w_i} \mathcal{L}$ and $\nabla_{w_i}^2 \mathcal{L}$.

Problem 2 (Gaussian kernels) – 15 points

We saw how we can construct a kernel between two points $u, v \in \mathbb{R}^d$ using the dot product (or integral) of their high-dimensional mappings $\phi(u)$ and $\phi(v)$. In the integral case,

$$k(u, v) = \int_{\mathbb{R}^d} \phi_t(u) \phi_t(v) dt,$$

where t is some parameter that is integrated out. Show that the mapping

$$\phi_t(u) = \frac{1}{(2\pi\beta')^{d/2}} \exp \left\{ -\frac{\|u - t\|^2}{2\beta'} \right\}$$

reproduces the Gaussian kernel $k(u, v) = \alpha \exp \left\{ -\frac{\|u - v\|^2}{\beta} \right\}$ for an appropriate setting of α and β .

Hint: This will be very difficult to do without using properties of multivariate Gaussians and their marginal distributions to draw some necessary conclusions. Try framing this as a probability question.

Problem 3 (Classification) – 70 points

In this problem you will implement three classifiers and run them on the MNIST Handwritten Digits data set posted on Courseworks and the class website. Do not do preprocessing to the data other than what is indicated at the end of the README below. The three classifiers must be implemented by you to receive full credit. Information about the data is given at the end of the assignment.

All three sub-problems ask for you to show your results in a 10×10 “confusion matrix” (call it C). This can be done as follows: For each of the 500 predictions you make from the test set, let y_t be the *true* class label and y_p be the *predicted* class label using your algorithm. Update $C(y_t, y_p) \leftarrow C(y_t, y_p) + 1$ for each prediction. At the end, C should sum to 500 and each row should sum to 50. (C can then be normalized, but leave it unnormalized for this assignment.)

Problem 3a (15 points) :

- Implement the k -NN classifier for $k = 1, 2, 3, 4, 5$.
- For each k calculate the confusion matrix and show the trace of this matrix divided by 500. This is the prediction accuracy. You don’t need to show the confusion matrix.
- For $k = 1, 3, 5$, show three misclassified examples and indicate the true class and the predicted class for each one (see the README below).

Problem 3b (25 points) :

- Implement the Bayes classifier using multivariate Gaussian distributions as the generative distribution for the data in each class.
- Derive the maximum likelihood estimate for the 10-dimensional distribution on classes and the Gaussian parameters for a particular class j that you will need for this problem.
- Show the confusion matrix in a table. As in Problem 3a, indicate the prediction accuracy by summing along the diagonal and dividing by 500.
- Show the mean of each Gaussian as an image using the provided Q matrix (see the README).
- Show three misclassified examples and show the probability distribution on the 10 digits learned by the Bayes classifier for each one.

Problem 3c (30 points) :

- Implement the multiclass logistic regression classifier you derived in Problem 1. You only need to use $\nabla_w \mathcal{L}$ to satisfy the requirements of this problem. In this case, you might want try a stepsize on the order of $\rho = 0.1/5000$.
- For each cycle through w_0, \dots, w_9 , calculate \mathcal{L} (see Problem 1) and plot as a function of iteration. Run your algorithm for 1000 iterations.
- Show the confusion matrix in a table. Indicated the prediction accuracy by summing along the diagonal and dividing by 500.
- Show three misclassified examples and show the probability distribution on the 10 digits learned by the softmax function for each one.

Option for Problem 3c : If you are unable to get the multiclass logistic regression classifier working (e.g., because of an issue with your derivation), you can instead implement a one-vs-rest classifier using the basic logistic regression model derived in class, then use this model to answer all questions in this problem. To get full credit, you will need to completely motivate the “one-vs-rest” classification setup and describe in detail how you used this framework to answer each question in this part of the assignment.

DATA README

The following data is included:

Xtrain : 20 x 5000 matrix of training data for digits 0 through 9.

label_train : 5000 dimensional vector containing the corresponding class (digit) in Xtrain.

Xtest : 20 x 500 matrix of testing data for digits 0 through 9.

label_test : 500 dimensional vector containing the corresponding class (digit) in Xtest.

Q : 784 x 20 matrix of principal components. You will only need this matrix for visualizing the images asked for. The original data contains 28 x 28 images that are vectorized into 784 x 1 vectors. I projected the data onto its first 20 principle components to reduce the dimensionality of the problem. These components are contained in Q. You can see the information captured by these 20 vectors as follows: For a 20 x 1 vector "x" from a training or testing matrix, let $y = Q \cdot x$. Then show this as an image, e.g., in Matlab using `imagesc(reshape(y,28,28))`. The image will probably need to be rotated and/or flipped. Some information is lost in this process (and the data was centered) so the resulting images won't be crystal clear, but the digit should be recognizable.

IMPORTANT (FOR PART 3c ONLY)

Before training and testing the classifier, be sure to add a 21st dimension to the data set equal to one in order to allow for a shift in the decision boundary! Therefore, the classification vectors w_0, \dots, w_9 should be 21 dimensions. And the 21st row of X show be entirely equal to one.